

RELATIONAL ALGEBRA

Dr. Mohammed A. Mohammed
University of Sulaimani | College of science | Computer dept.

Relational Query Languages

- ***Query languages:*** Allow manipulation and retrieval of data from a database.
- Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- **Query Languages \neq programming languages!**
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

Formal Relational Query Languages

Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:

- **Relational Algebra:** More **operational**, very useful for representing execution plans.
- **Relational Calculus:** Lets users describe what they want, rather than how to compute it. (**Non-operational, declarative.**)

Preliminaries (Introductions)

A query is applied to *relation instances*, and the result of a query is also a relation instance.

```
SELECT S.name, E.grade  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid and S.age<20
```

- Each relational query describes a **step-by-step procedure** for computing the desired answer, based on the order in which operators are applied in the query.
- The procedural nature of the algebra allows us to think of an algebra expression as a guidelines, or a plan, for evaluating a query, and relational systems in fact use algebra expressions to represent query evaluation plans.

Notation

Positional vs. named-field notation:

- **Positional notation** – referring to fields by their positions. Easier for formal definitions.
- **Named-field notation** – using field names to refer to fields. Make query more readable.
- Both used in SQL

Sample Schema

- We present a number of sample queries using the following schema:

Sailors(*sid: integer*, *sname: string*, *rating: integer*, *age: real*)

Boats(*bid: integer*, *bname: string*, *color: string*)

Reserves(*sid: integer*, *bid: integer*, *day: date*)

Example Instances

- “Sailors” and “Reserves” relations for our examples.
- We’ll use positional or named field notation, assume that names of fields in query results are ‘inherited’ from names of fields in query input relations.

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Relational Algebra

- Basic operations:

- Selection (σ) Selects a subset of **rows** from relation.
- Projection (π) Deletes unwanted **columns** from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

- Additional operations:

- Intersection, join, division, renaming: Not essential, but (very!) useful.

- Since each operation returns a relation, operations can be *composed*!

Projection

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

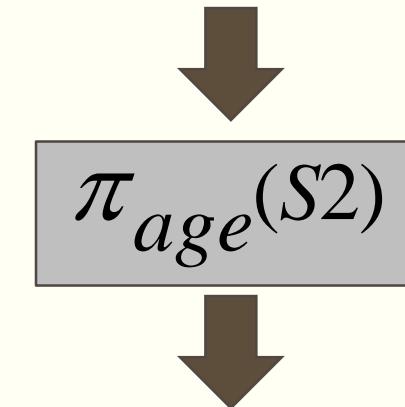
sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

Projection

- Deletes attributes that are not in *projection list*.
- *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the (only) input relation.
- Projection operator has to eliminate *duplicates*.
 - **Note:** real systems typically don't do duplicate elimination unless the user explicitly asks for it.

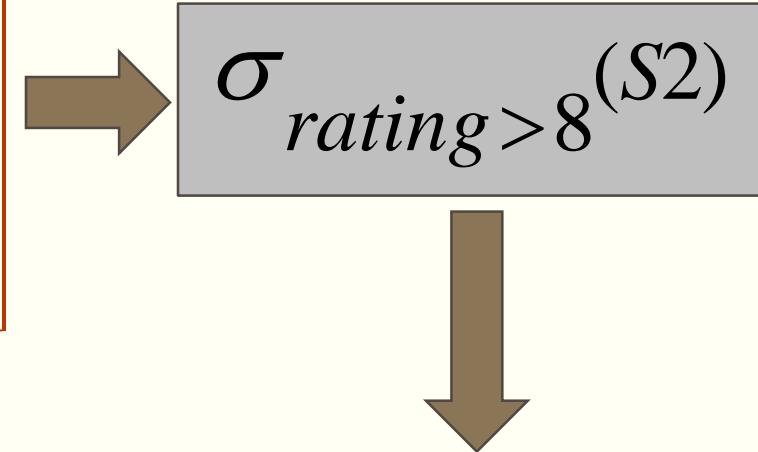
sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



age
35.0
55.5

Selection

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



- Selects rows that satisfy selection condition.
- *Schema* of result identical to schema of (only) input relation

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

Composition Operator

Result relation can be the input for another relational algebra operation!

sid	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$

sname	rating
yuppy	9
rusty	10

Union

s_1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

s_2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S_1 \cup S_2$

الناتج هو مجموعتين

Intersection

s_1	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

s_2	<u>sid</u>	sname	rating	age
	28	yuppy	9	35.0
	31	lubber	8	55.5
	44	guppy	5	35.0
	58	rusty	10	35.0

	<u>sid</u>	sname	rating	age
	31	lubber	8	55.5
	58	rusty	10	35.0

$S_1 \cap S_2$
↓
elimi... disc...

Set-Difference

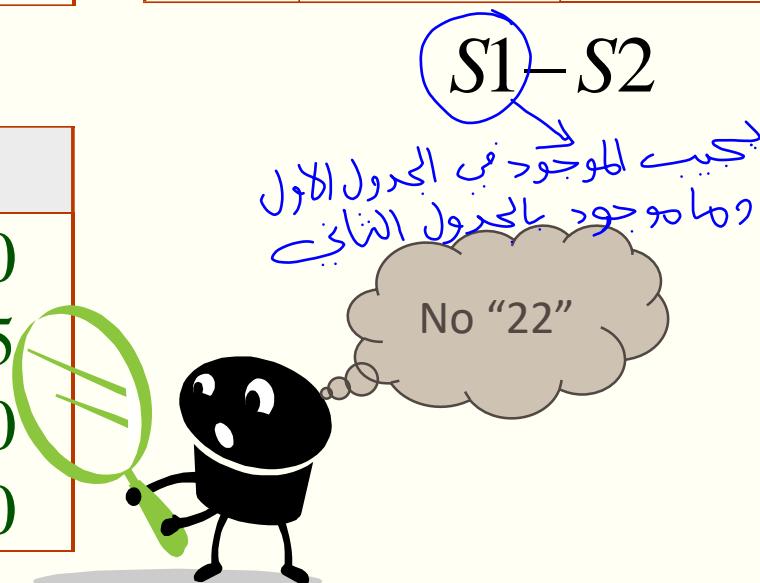
s_1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

sid	sname	rating	age
22	dustin	7	45.0

s_2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0



Union, Intersection, Set-Difference

All of these operations take two input relations, which must be *union-compatible*:

- Same number of fields.
- ‘Corresponding’ fields have the same type.

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$$S1 \cup S2$$

sid	sname	rating	age
22	dustin	7	45.0

$$S1 - S2$$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$$S1 \cap S2$$

Cross-Product

<i>s1</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>R1</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

$R1 \times S1$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Each row of R1 is paired with each row of S1.

Cross-Product

Result schema has one field per field of S1 and R1, with field names ‘inherited’ if possible.

- Conflict: Both S1 and R1 have a field called *sid*.

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Renaming operator:

$$\rho(C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$$

ghost class
rename cycle
↓ will go to

C(*sid1*:integer, *sname*: string, *rating*: integer, *age*: real, *sid2*: integer, *bid*: integer, *day*: dates)

Joins

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Sailors

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Boats

<u>bid</u>	<u>bname</u>	<u>color</u>
101	Interlake	Blue
102	Interlake	Red
103	Clipper	Green
104	Marine	Red

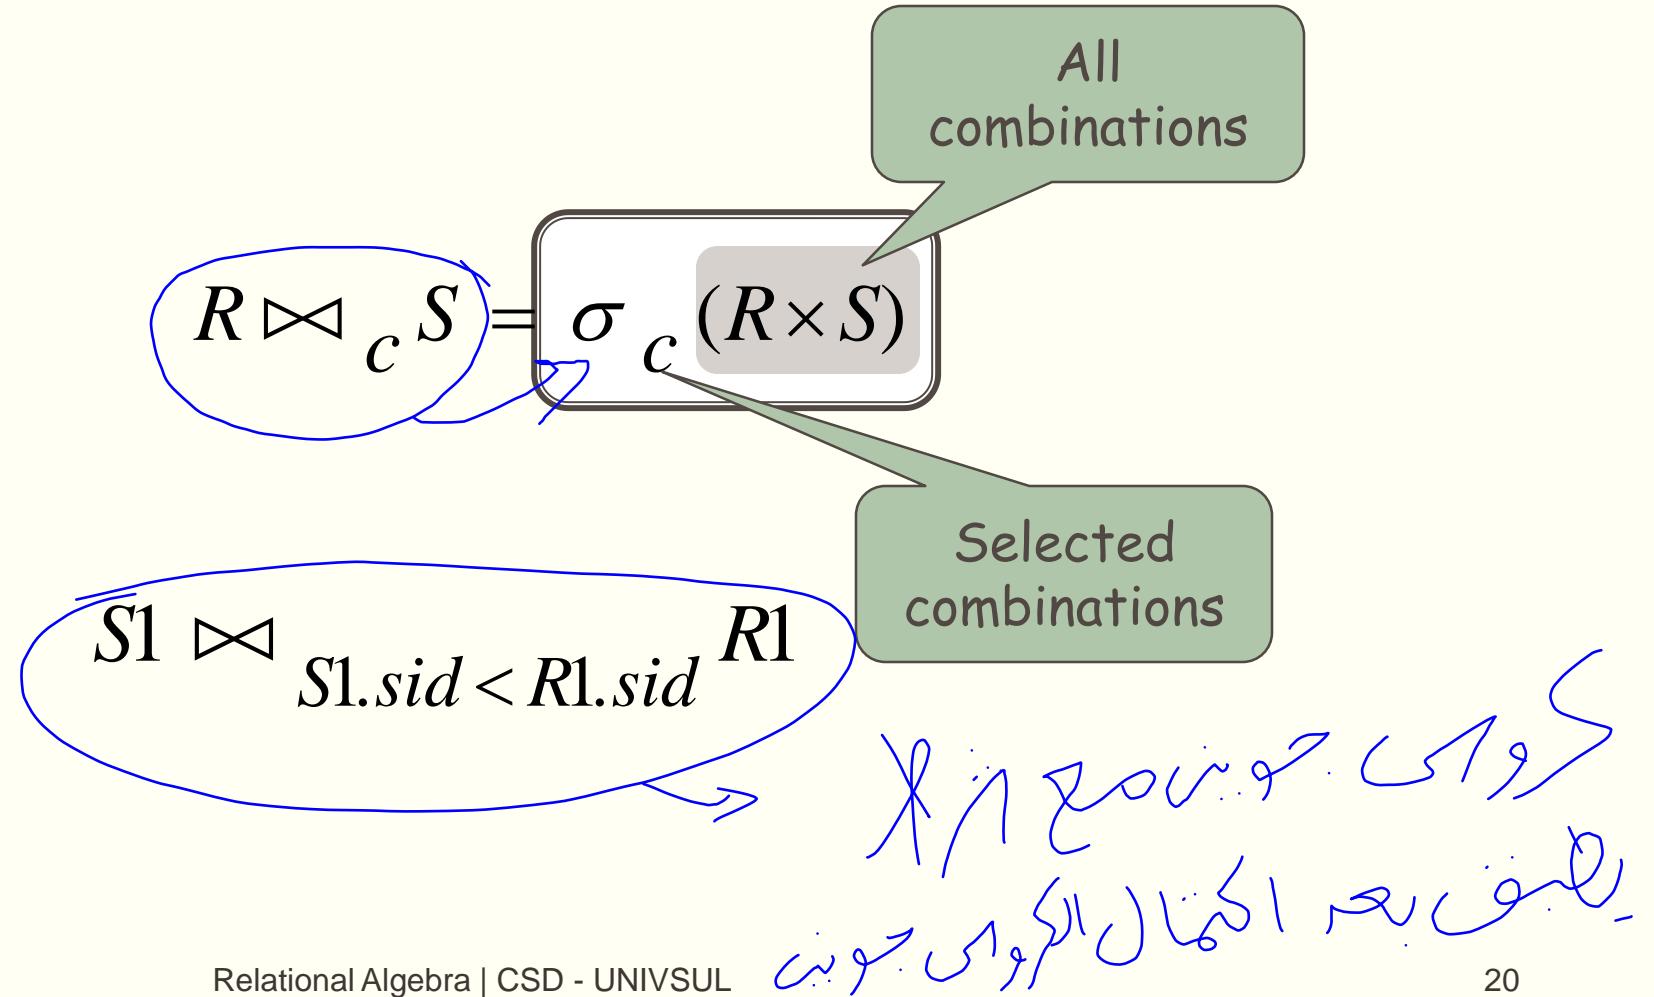
Find names of sailors who've reserved a green boat

($\pi_{sname} ((\sigma_{color='Green'} Boats) \bowtie Reserves) \bowtie Sailors)$)

Joins

Condition Join:

Example:



Joins

$$R \bowtie_c S = \sigma_c (R \times S)$$

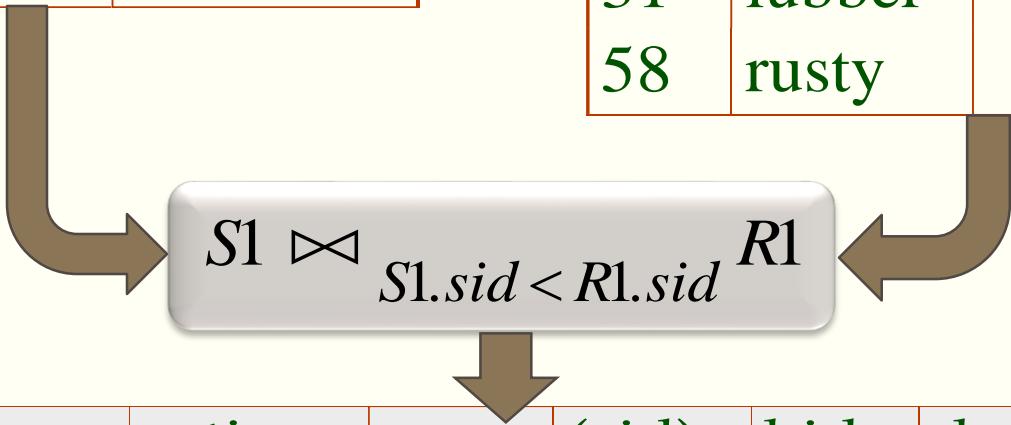
Condition Join:

$R1$

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

$S1$

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0


$$S1 \bowtie_{S1.sid < R1.sid} R1$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

Joins

$$S1 \bowtie_{S1.sid = R1.sid} R1$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

$$R1 \times S1$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Joins

- *Result schema* same as that of cross-product.
- Fewer tuples than cross-product, might be able to compute more efficiently
- Sometimes called a *theta-join*.

Equi-Join

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	day
22	101	10/10/96
58	103	11/12/96

Foreign key

Equi-Join: A special case of condition join where the condition c contains only **equalities**.

$S1 \bowtie_{sid} R1 =$

جوجو جوجو
22 101, 58 103

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

Equi-Joins & Natural Join

- **Equi-Join:** A special case of condition join where the condition c contains only *equalities*.

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

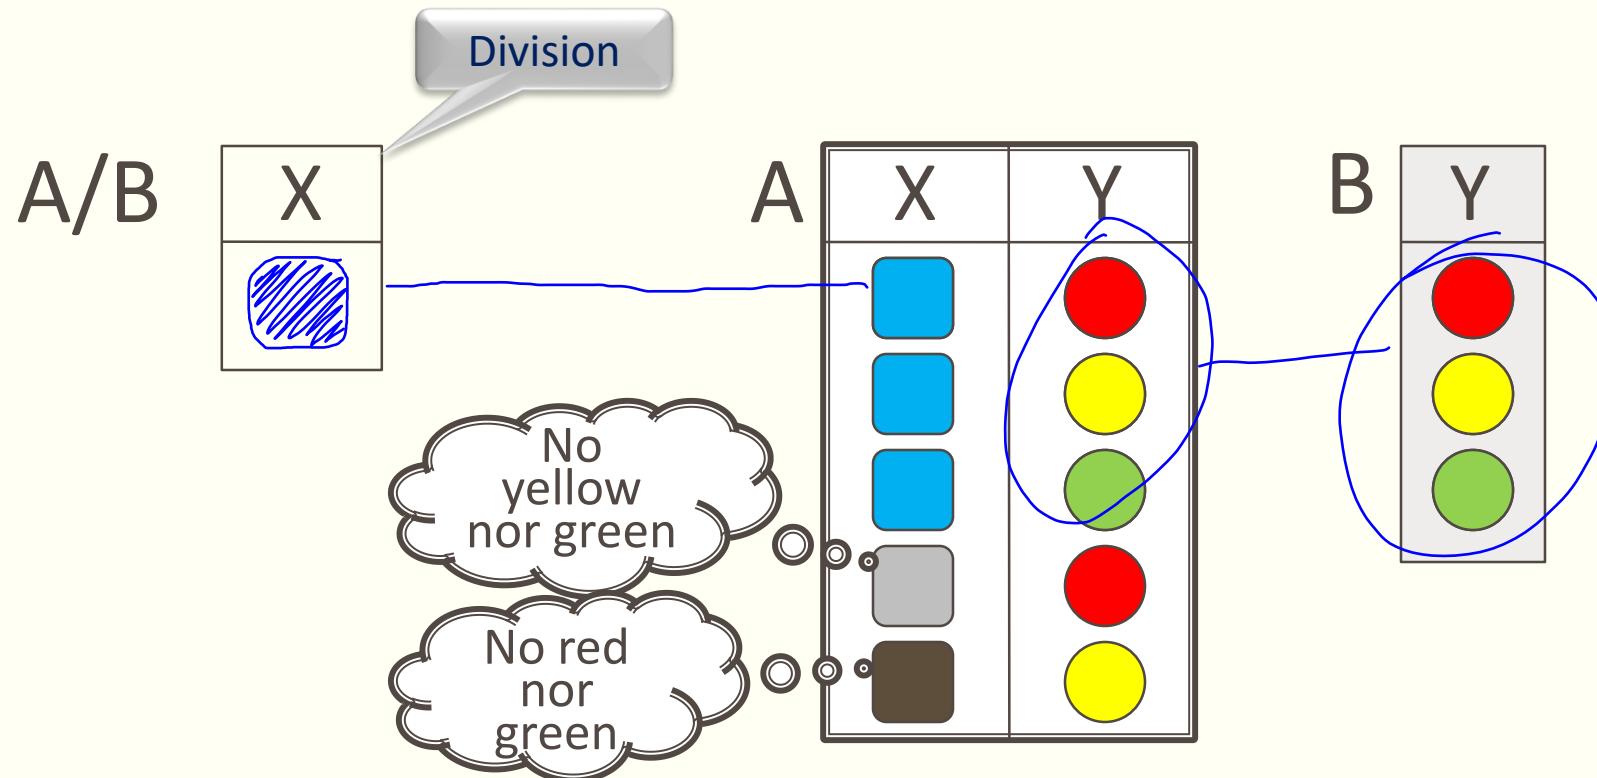
$$S1 \bowtie_{sid} R1$$

- Result schema similar to cross-product, but **only one copy of fields** for which equality is specified.
- **Natural Join:** Equijoin on *all* common fields.

Division

Not supported as a primitive operator, but useful for expressing queries like:

Find sailors who have reserved all boats



Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

B1

sno
s1
s2
s3
s4

A/B1

pno
p2
p4

B2

sno
s1
s4

A/B2

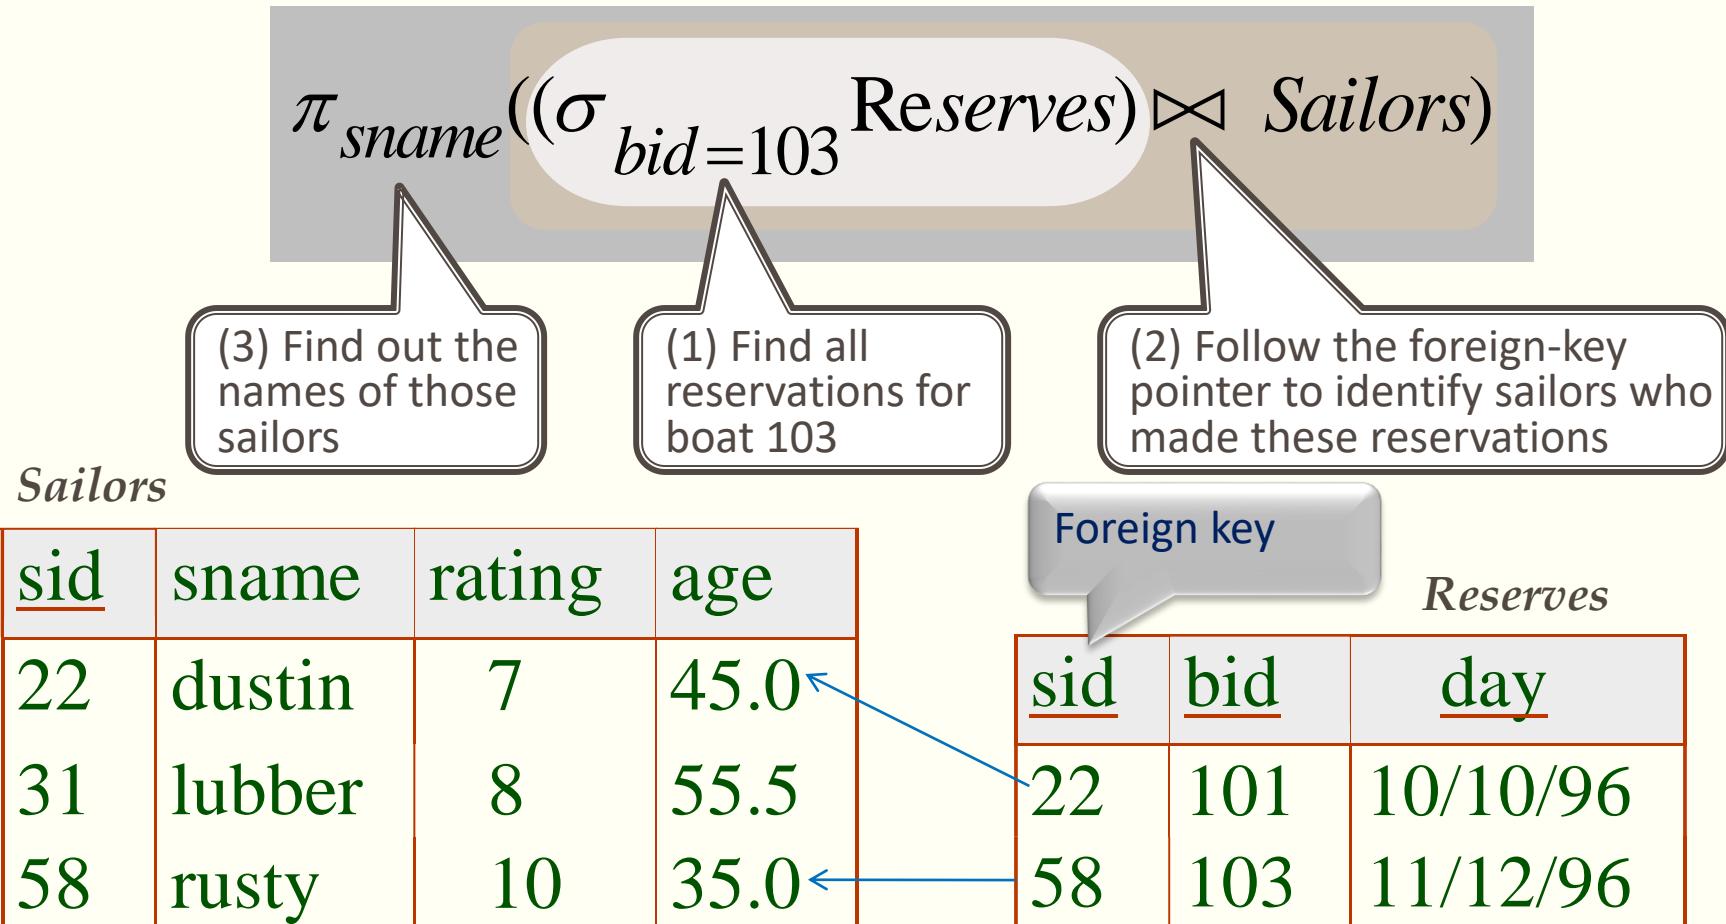
pno
p1
p2
p4

B3

sno
s1

A/B3

Find names of sailors who've reserved boat #103



Find names of sailors who've reserved boat #103

Solution 2:

$$\rho (\text{Temp1}, \sigma_{bid=103} \text{Reserves})$$

(1) Find all reservations for boat 103

$$\rho (\text{Temp2}, \text{Temp1} \bowtie \text{Sailors})$$

(2) Follow the foreign-key pointer to identify sailors who made these reservations

Sailors $\pi_{sname}(\text{Temp2})$

(3) Find out the names of those sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Foreign key

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Find names of sailors who've reserved boat #103

Solution 3:

$$\pi_{sname}(\sigma_{bid=103}(\text{Reserves} \bowtie \text{Sailors}))$$

(3) Find out the names of sailors who made those sailors

(2) Retain only reservations for boat 103

(1) Find reservations for every sailor

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Foreign key

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Which one is better ? 💬


$$\pi_{sname}((\sigma_{bid=103} \text{Reserves}) \bowtie \text{Sailors})$$

bid₁₀₃ op op Jg1
equal join Sgn, op Jg1


$$\pi_{sname}(\sigma_{bid=103} \underline{(\text{Reserves} \bowtie \text{Sailors})})$$

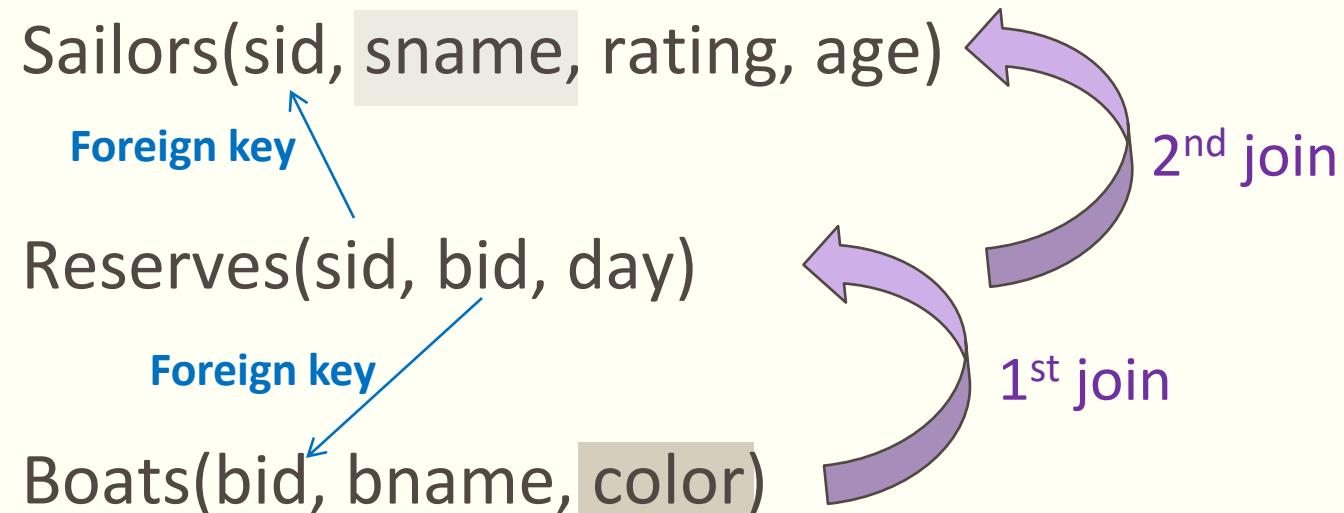
equal join Sgn, op Jg1

Find names of sailors who've reserved a red boat

Start with "boat"
because we have a
predicate about boat

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

1 2



Find names of sailors who've reserved a red boat

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

- ❖ A more efficient solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} Boats) \bowtie Res) \bowtie Sailors)$$


Make Join less expensive

A query optimizer can find this, given the first solution!

Find sailors who've reserved a red or a green boat

- Can identify all red or green boats, then find sailors who've reserved one of these boats:

$$\rho (Tempboats, (\sigma_{color='red' \vee color='green'} Boats))$$
$$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$$

- Can also define Tempboats using union

$$\rho (Tempboats, (\sigma_{color='red'} Boats) \cup (\sigma_{color='green'} Boats))$$

Find sailors who've reserved a red and a green boat

This won't work!

Empty Relation



$\rho(\text{Tempboats}, (\sigma_{\text{color}=\text{'red'}} \wedge \sigma_{\text{color}=\text{'green'}}) \text{ Boats})$

$\pi_{\text{sname}}(\text{Tempboats} \bowtie \text{Reserves} \bowtie \text{Sailors})$

Find sailors who've reserved a red and a green boat

- Identify sailors who've reserved red boats

$$\rho (Tempred, \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$$

- Identify sailors who've reserved green boats

$$\rho (Tempgreen, \pi_{sid}((\sigma_{color='green'} Boats) \bowtie Reserves))$$

- Find the intersection

$$\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$$

Find the names of sailors who've reserved all boats

- Uses division; schemas of the input relations to division must be carefully chosen:

$$\rho (Tempsids, (\pi_{sid,bid} \text{Reserves}) / (\pi_{bid} \text{Boats}))$$

All boats

$$\pi_{sname} (Tempsids \bowtie Sailors)$$

- To find sailors who've reserved all 'Interlake' boats:

.....

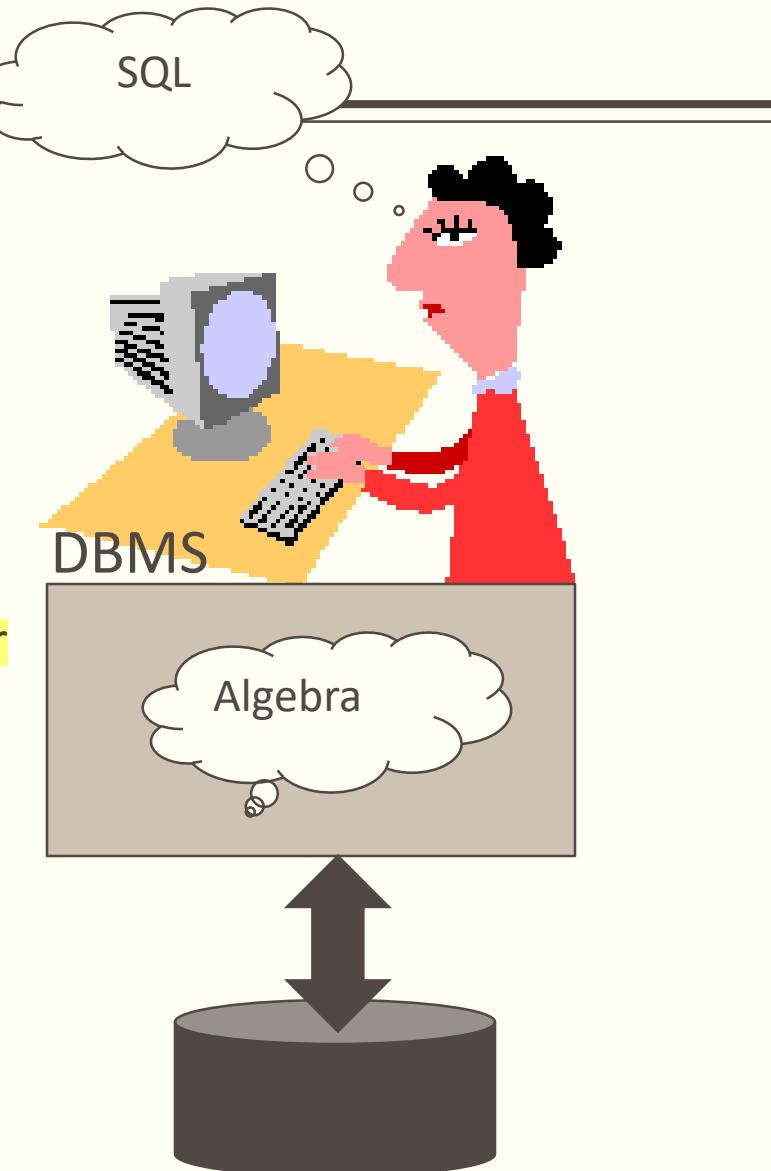
$$/ \pi_{bid} (\sigma_{bname='Interlake'} \text{Boats})$$

'Interlake' boats

Operation	Purpose	Notation
SELECT	Selects all tuples that satisfy the selection condition from a relation R .	$\sigma_{\text{selection condition}}(R)$
PROJECT	Produces a new relation with only some of the attributes of R , and removes duplicate tuples.	$\pi_{\text{attribute list}}(R)$
THETA JOIN	Produces all combinations of tuples from R_1 and R_2 that satisfy the join condition.	$R_1 \bowtie_{\text{join condition}} R_2$
EQUIJOIN	Produces all the combinations of tuples from R_1 and R_2 that satisfy a join condition with only equality comparisons.	$R_1 \bowtie_{\text{join condition}} R_2,$ OR $R_1 \bowtie_{(<\text{join attributes 1}>), (<\text{join attributes 2}>)} R_2$
NATURAL JOIN	Same as EQUIJOIN except that the join attributes of R_2 are not included in the resulting relation; if the join attributes have the same names, they do not have to be specified at all.	$R_1 *_{\text{join condition}} R_2,$ OR $R_1 *_{(<\text{join attributes 1}>), (<\text{join attributes 2}>)} R_2$ OR $R_1 * R_2$
UNION	Produces a relation that includes all the tuples in R_1 or R_2 or both R_1 and R_2 ; R_1 and R_2 must be union compatible.	$R_1 \cup R_2$
INTERSECTION	Produces a relation that includes all the tuples in both R_1 and R_2 ; R_1 and R_2 must be union compatible.	$R_1 \cap R_2$
DIFFERENCE	Produces a relation that includes all the tuples in R_1 that are not in R_2 ; R_1 and R_2 must be union compatible.	$R_1 - R_2$
CARTESIAN PRODUCT	Produces a relation that has the attributes of R_1 and R_2 and includes as tuples all possible combinations of tuples from R_1 and R_2 .	$R_1 \times R_2$

Summary

- The relational model has rigorously defined query languages that are simple and powerful.
- Relational algebra is more operational; useful as internal representation for query evaluation plans.
- Several ways of expressing a given query; a query optimizer should choose the most efficient version.



Any Question ?

